

The Kernel Report

Namhyung Kim
LG Electronics

Kernel versions

- **4.16**

- HRTimer in softirq, printk() lockup prevention, Usercopy whitelisting

- **4.17**

- Inter-event tracing histogram, scheduler load tracking improvements (+ 4.19)

- **4.18**

- Restartable sequence, bpfILTER, dm-writecache


- **4.19**

- Async IO polling, block IO latency controller, L1TF

HRTimer in softirq

- **Running (high-resolution) timer function in a soft-irq context instead of hard-irq**
 - `HRTIMER_MODE_{ABS,REL}_SOFT`
- **Hard-irq vs Soft-irq**
 - Hardirq handler runs with irq-disabled
 - Most of work deferred to the bottom-half
 - softirq, tasklet or workqueue
 - Or threaded-irq handler (rt-task)

Interrupt handling



printk() lockup prevention


- **Lockup senario**

- First printk() will output all contents of logbuf
- A burst of printk() will soft-lockup the first caller if it's faster than console ouput

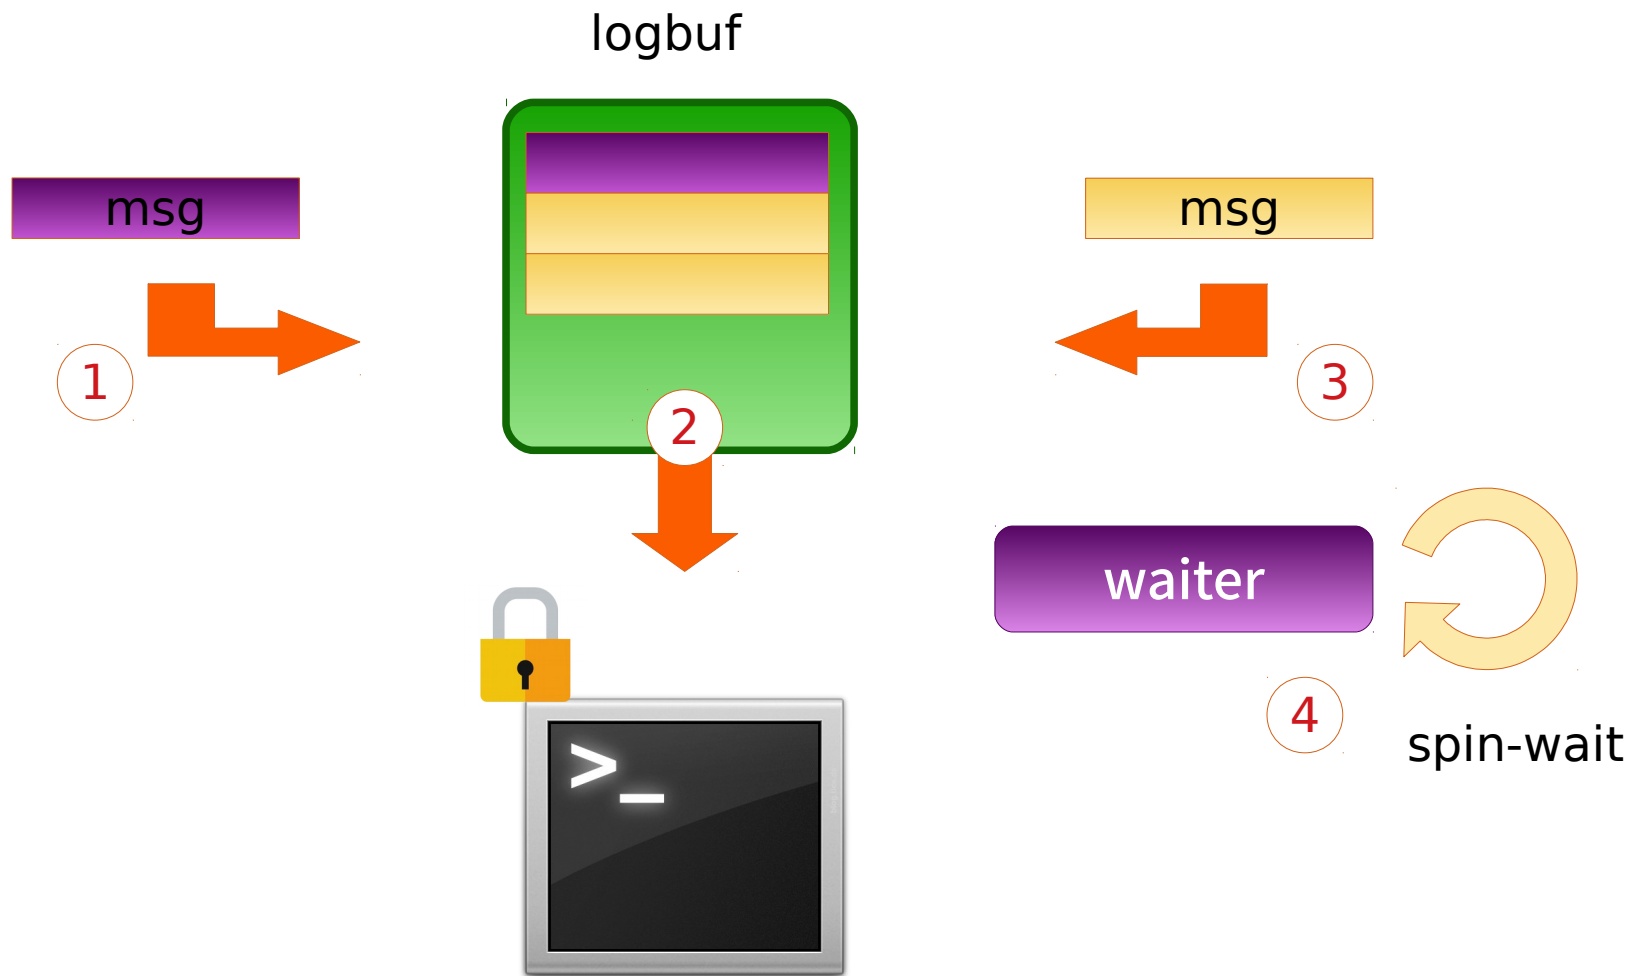
- **Solution**

- New code will change responsibility of output to next caller

printk: before



printk: After



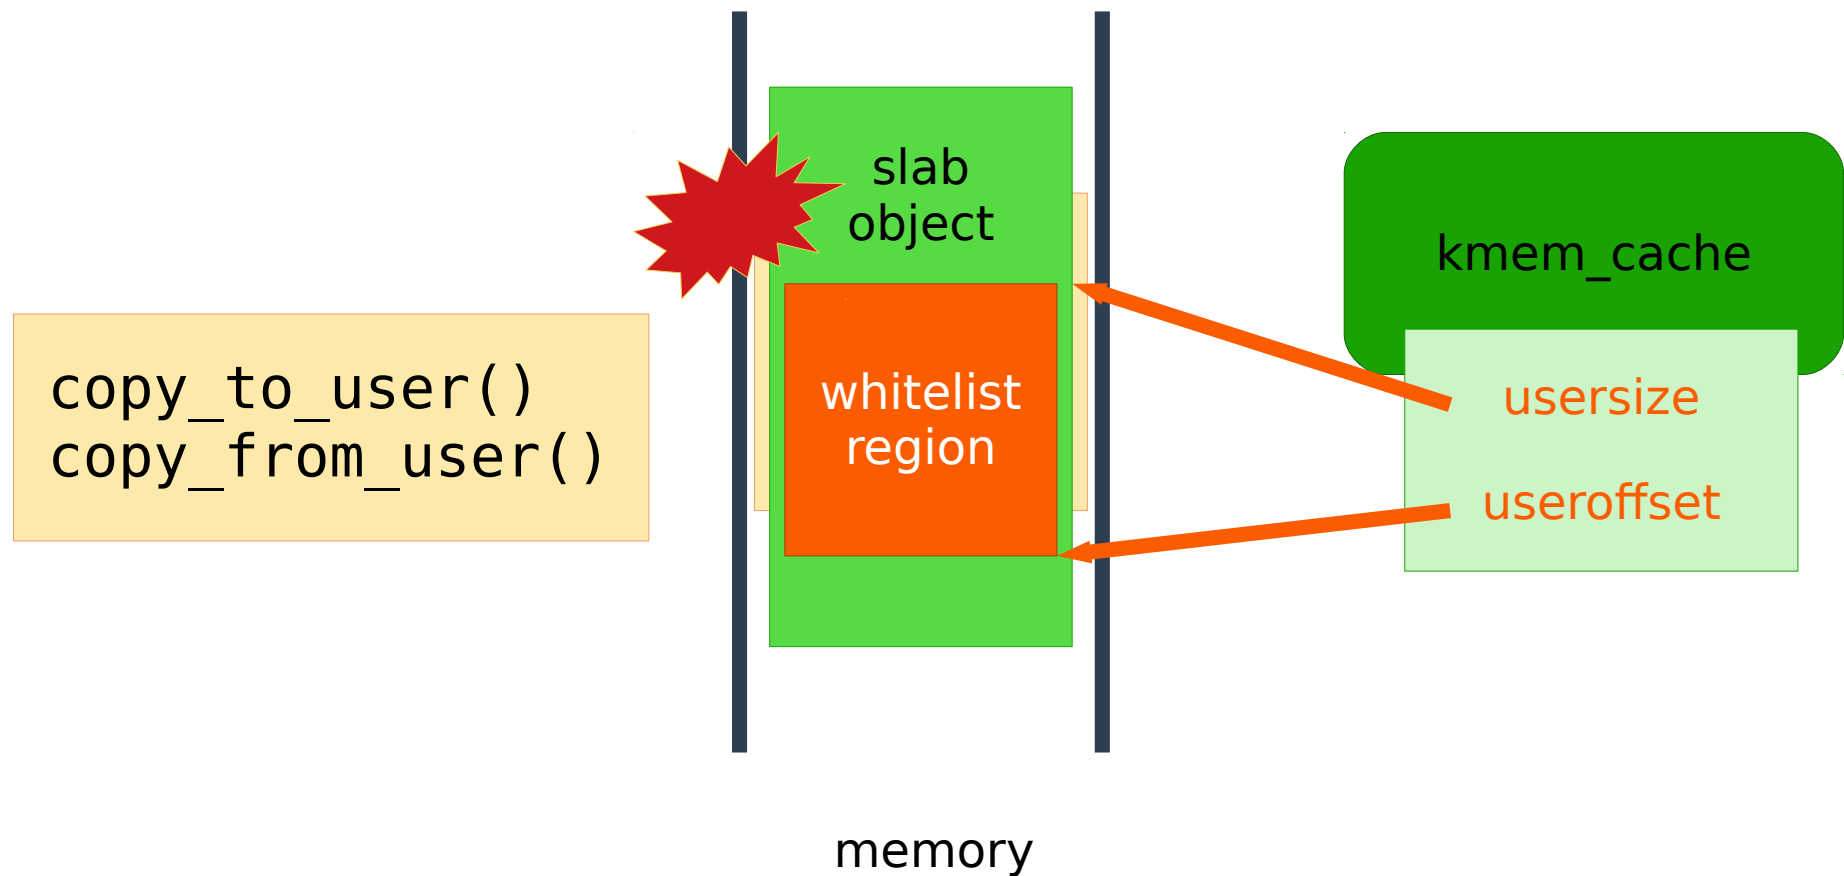
Usercopy whitelisting

- **Security feature**

- Limit range of interaction between user and kernel in a slab object

```
struct kmem_cache *  
kmem_cache_create_usercopy(const char *name,  
                           unsigned int size, unsigned int align,  
                           slab_flags_t flags,  
                           unsigned int useroffset, unsigned int usersize,  
                           void (*ctor)(void *));
```


Usercopy whitelisting



Scheduler load tracking

- **PELT**
 - Per-Entity Load Tracking
 - Keep (cpu) util for each task (sched-entity)
- **UTIL_EST (4.17)**
 - Save util value when task goes to sleep
- **IRQ/Real-time utilization (4.19)**
 - For schedutil cpufreq governor

Load Tracking

- **Load = CPU utilization**
- **Task util**
 - RUNNING / time (moving average)
 - Task placement decision
- **CPU util**
 - Sum of task utils
 - Load balancing decision
 - cpufreq decision

Util Estimation

- **UTIL_EST**
 - Enqueued : task util at the time of dequeue
 - Ewma : exp. Weighted moving average of 'enqueued'. (task only)
- **Use max value of task util**
 - Insensitive to transient changes
 - Better estimate a long slept big task
 - FAIR tasks only

Inter-event tracing histogram

- **Inter-event tracing**

- Tracing two or more events and calculate value from those events
- Create new (synthesized) event

- **Tracing histogram**

- don't save all trace record
- Aggregate result using given keys

Tracing histogram

- **To know distribution of values**
 - Just count the event (using keys)
 - don't waste buffer and no need to copy

* **kmalloc event histogram**

Key: pid	hitcount	Value: bytes_alloc
1234	1	128
5678	3	768
13579	100	6400

Inter-event tracing

- **Inter-event tracing**

- Variable support : save last value
- Trigger action: do something if event match
- Synthesized event: create new event runtime

- **Wakeup latency**

- Time between task wake and schedule
- Using inter-event tracing
- Histogram on the synthesized event

Wakeup Latency example

```
# cd /sys/kernel/debug/tracing

# echo 'wakeup_latency u64 lat; pid_t pid; int prio' \
  >> synthetic_events

# echo 'hist:key=lat.log2' \
  >> events/synthetic/wakeup_latency/trigger

# echo 'hist:key=pid:t0=common_timestamp.usecs if comm=="cyclicttest"' \
  >> events/sched/sched_waking/trigger

# echo 'hist:key=next_pid:lat=common_timestamp.usecs-$t0: \
  onmatch(sched.sched_waking). \
  wakeup_latency($lat,next_pid,next_prio) \
  if next_comm=="cyclicttest" \
  >> events/sched/sched_switch/trigger

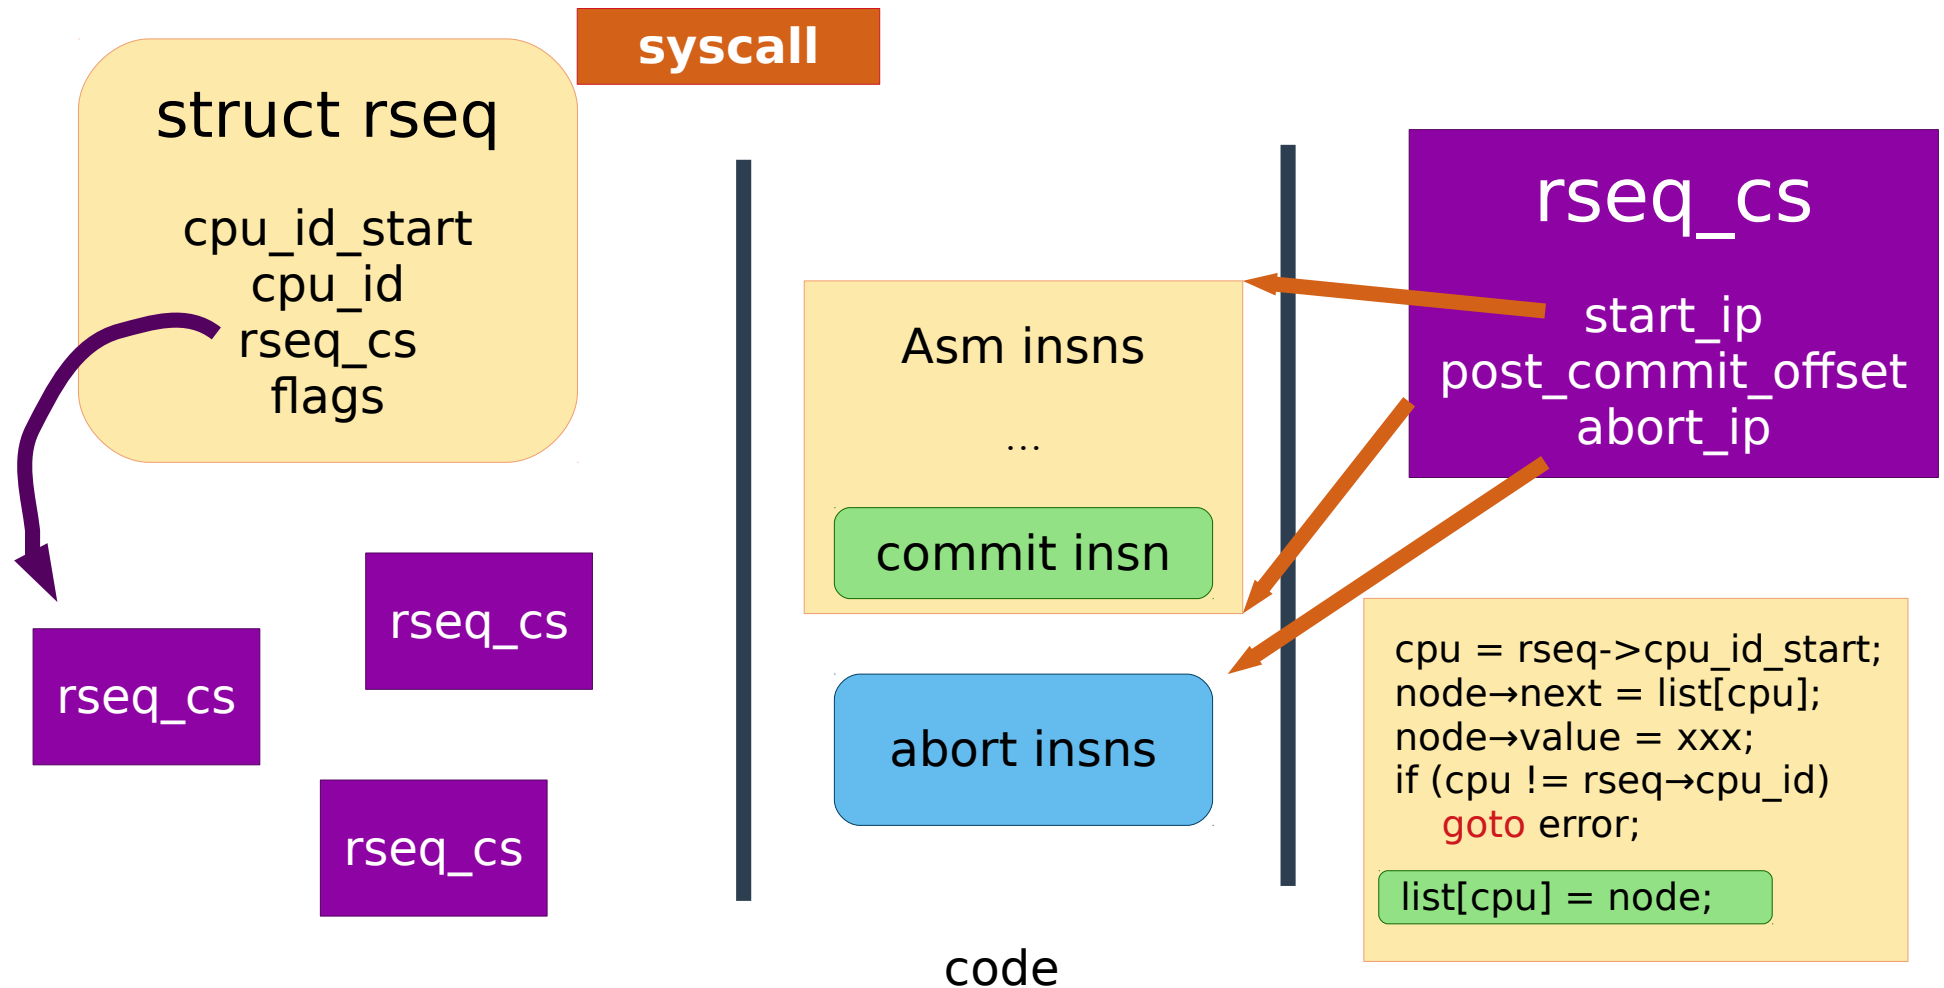
# cat events/synthetic/wakeup_latency/hist
```


Restartable Sequence

- **Fast Lockless Synchronization**
 - Make fast-path faster
 - Avoid atomic instruction
 - Speed up by using per-cpu data
- **New syscall**

```
long rseq(struct rseq *rseq, uint32_t len,  
          int flags, uint32_t sig);
```

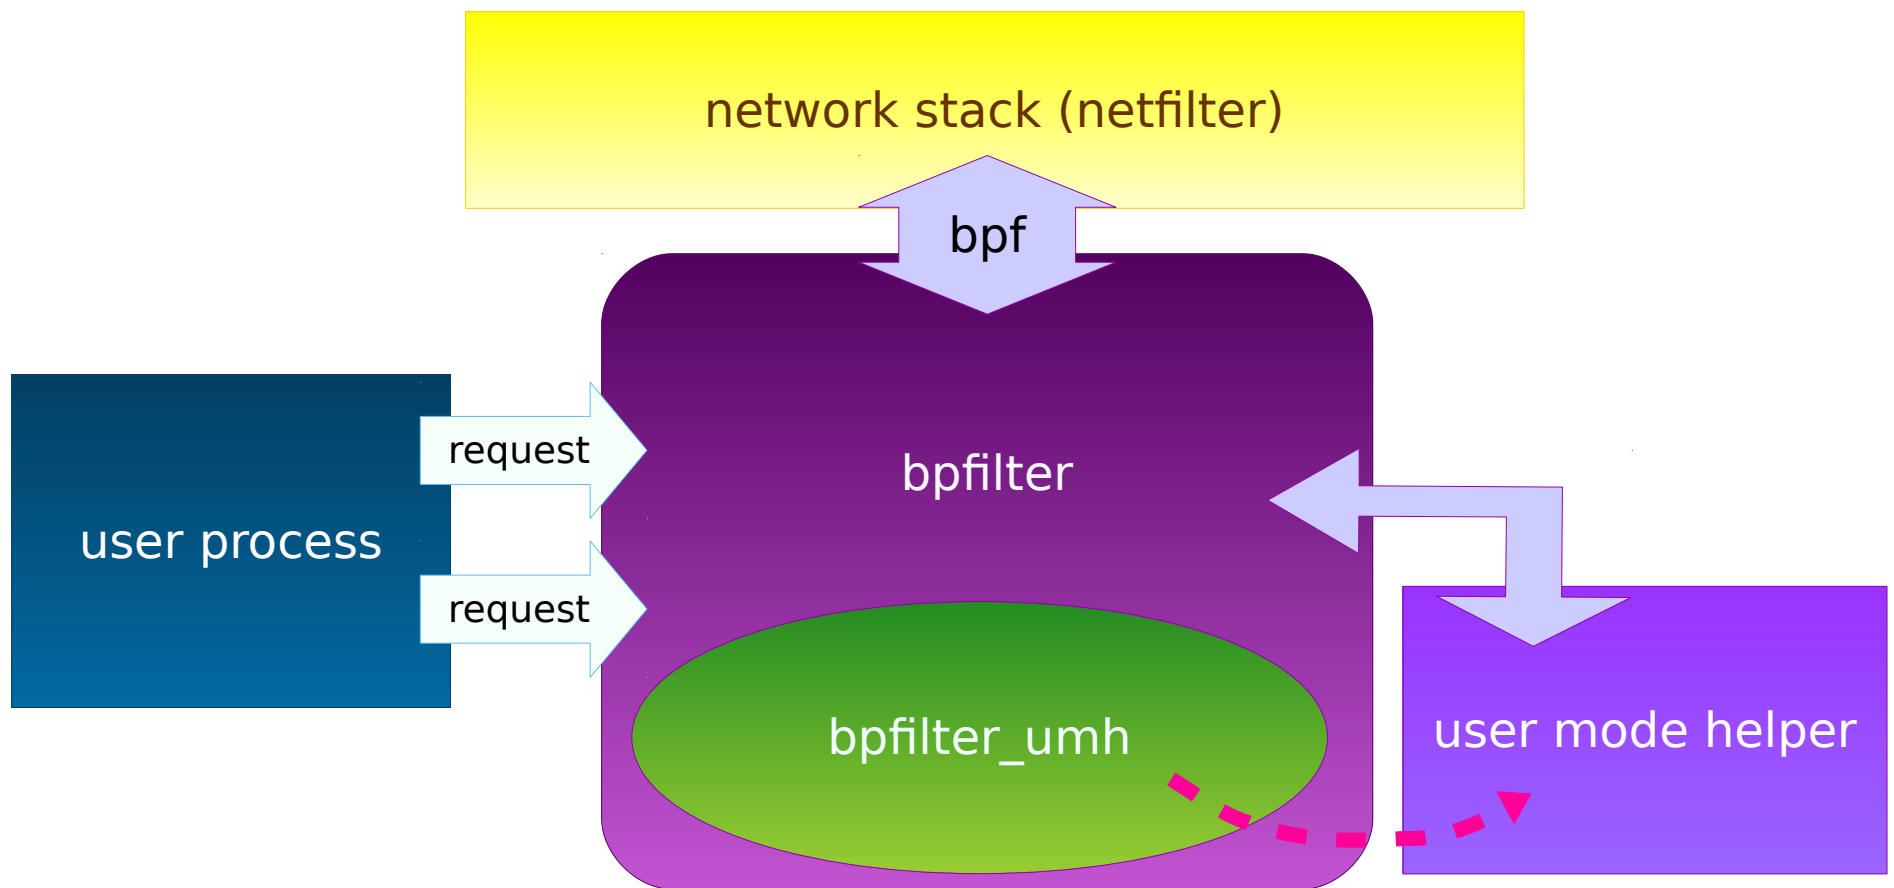
Restartable Sequence



bpfilter

- **Netfilter with BPF**
 - Replace iptables, nftables
 - Fast, safe vm in kernel
 - Initial work just started
- **User-mode helper**
 - Do complex jobs in user space
 - Kernel module include an ELF binary
 - Communicate via pipe

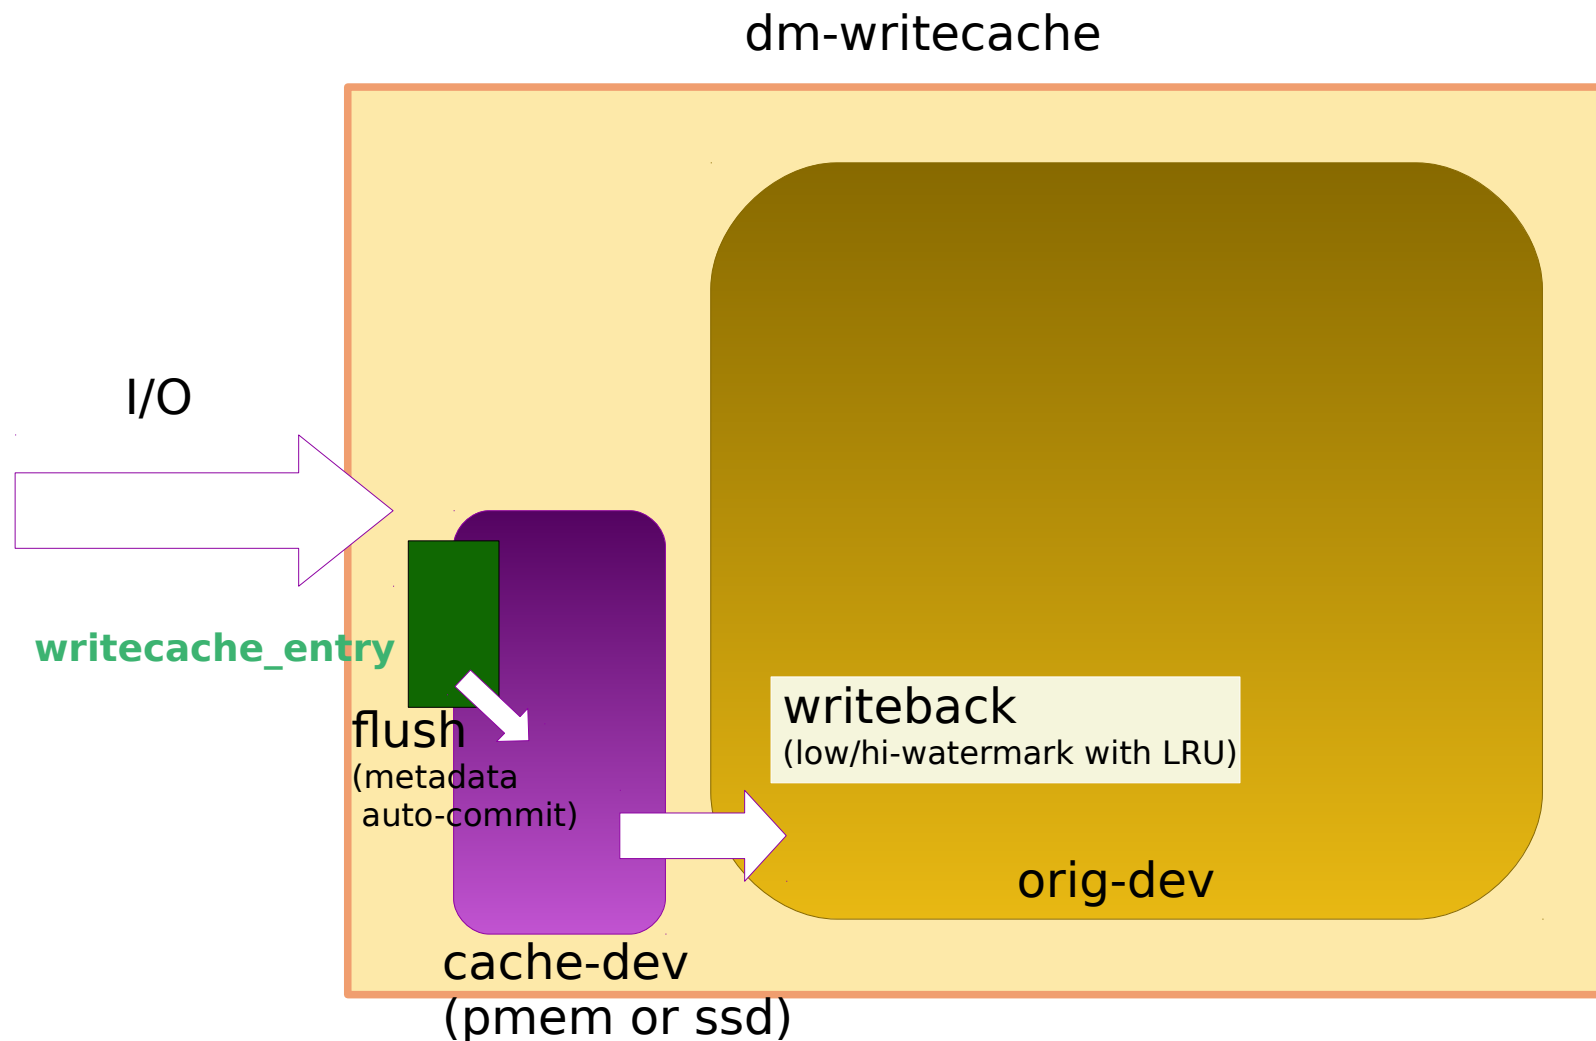
bpfilter



dm-writecache

- **Device mapper target**
 - Writeback caching to Pmem or SSD
 - Read will use page-cache (in RAM)
- **Operations**
 - Flush (commit) – metadata sync
 - Writeback : save to original dev

dm-writecache



Async I/O polling

- **New polling API**

- select()
- poll()
- epoll_wait()
- io_submit() (opcode = IOCB_CMD_POLL)

- **Kernel Async-IO API**

- No need to wait/block
- use ring-buffer for communication

Block I/O latency controller

- **Bandwidth vs Latency**

- Throttle I/O using cgroups

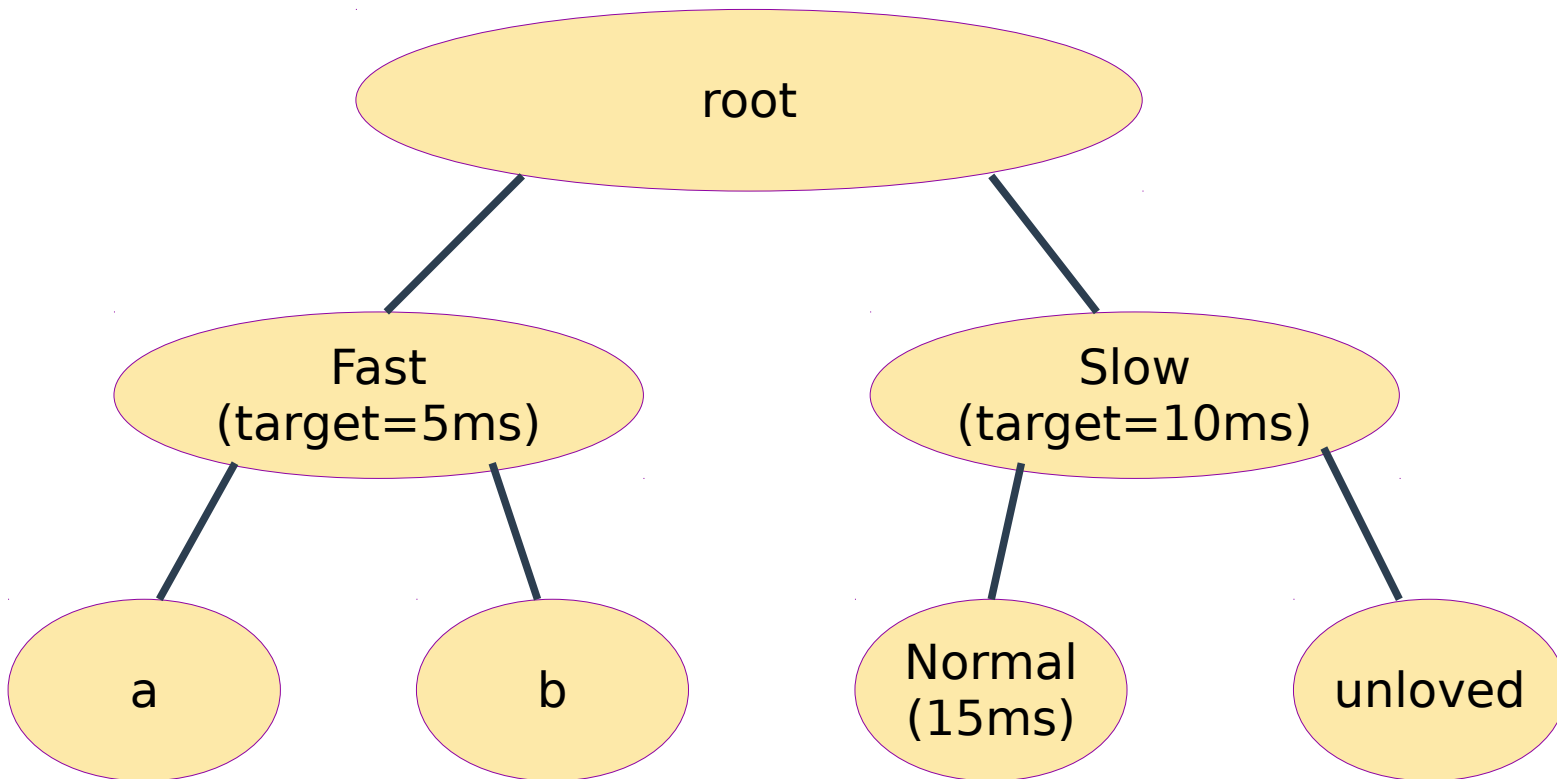
- **I/O latency controller**

- echo *<major:minor>* target=100 > io.latency
- *<major:minor>* = device ID
- target = *<max latency in usec>*

- **Exceptions**

- Metadata I/O (started by filesystem)
- Swap I/O (started by reclaim)

Block I/O latency controller



L1TF Mitigations

- **L1 Terminal Fault (Foreshadow)**
 - Affects Intel CPUs only
 - Allow access to any physical memory
- **Bypass “Present” bit in PTE**
 - Only in L1 cache with speculation



<https://lwn.net/Articles/762570/>

L1TF Mitigations

- **Mitigations**

- Invert all bits in PTE (for non-present page)
- Flush cache before return to user
- Disable SMT (hyper-threading)

- **Virtualization**

- Guest can run any kernel
- Disable EPT (Extended Page Table)



Any question?