

Kernel Report

v5.2 ~ v.5.3

Byungchul Park

byungchul.park@lge.com

max.byungchul.park@gmail.com

5.2

+ A clone(2) flag CLONE_PIDFD

Kill 121 !



+ A clone(2) flag CLONE_PIDFD



+ A clone(2) flag CLONE_PIDFD

121



+ A clone(2) flag CLONE_PIDFD



+ A clone(2) flag CLONE_PIDFD



+ A clone(2) flag CLONE_PIDFD

Kill 121 ?!



Background **PID reuse**



+ A clone(2) flag CLONE_PIDFD

Kill 121 !



+ A clone(2) flag CLONE_PIDFD

Kill 121 !



File Descriptor !

`open(2) ~ close(2)`

Background **pidfd_send_signal(2)**

A system call newly added in v5.1



+ A clone(2) flag **CLONE_PIDFD**

```
clone(..., CLONE_PIDFD, ..., &pidfd);
```

```
pidfd_send_signal(pidfd);
```

```
close(pidfd);
```

+ PSI notification



+ PSI notification



CPU

Memory

IO

+ PSI notification



CPU

Memory

IO

+ PSI notification



CPU

Memory

IO

Background PSI

Pressure Stall Information



CPU

Memory

IO

+ PSI notification



CPU

Memory

IO

+ PSI notification



CPU

Memory

IO

+ PSI notification

```
fds.fd = open("/proc/pressure/memory", flags);
```

```
write(fds.fd, "some 150000 1000000", len);
```

```
poll(&fds, 1, -1);
```


! BFQ improved



! BFQ improved



Background **IO scheduler**



Background **BFQ**

Budget Fair Queueing

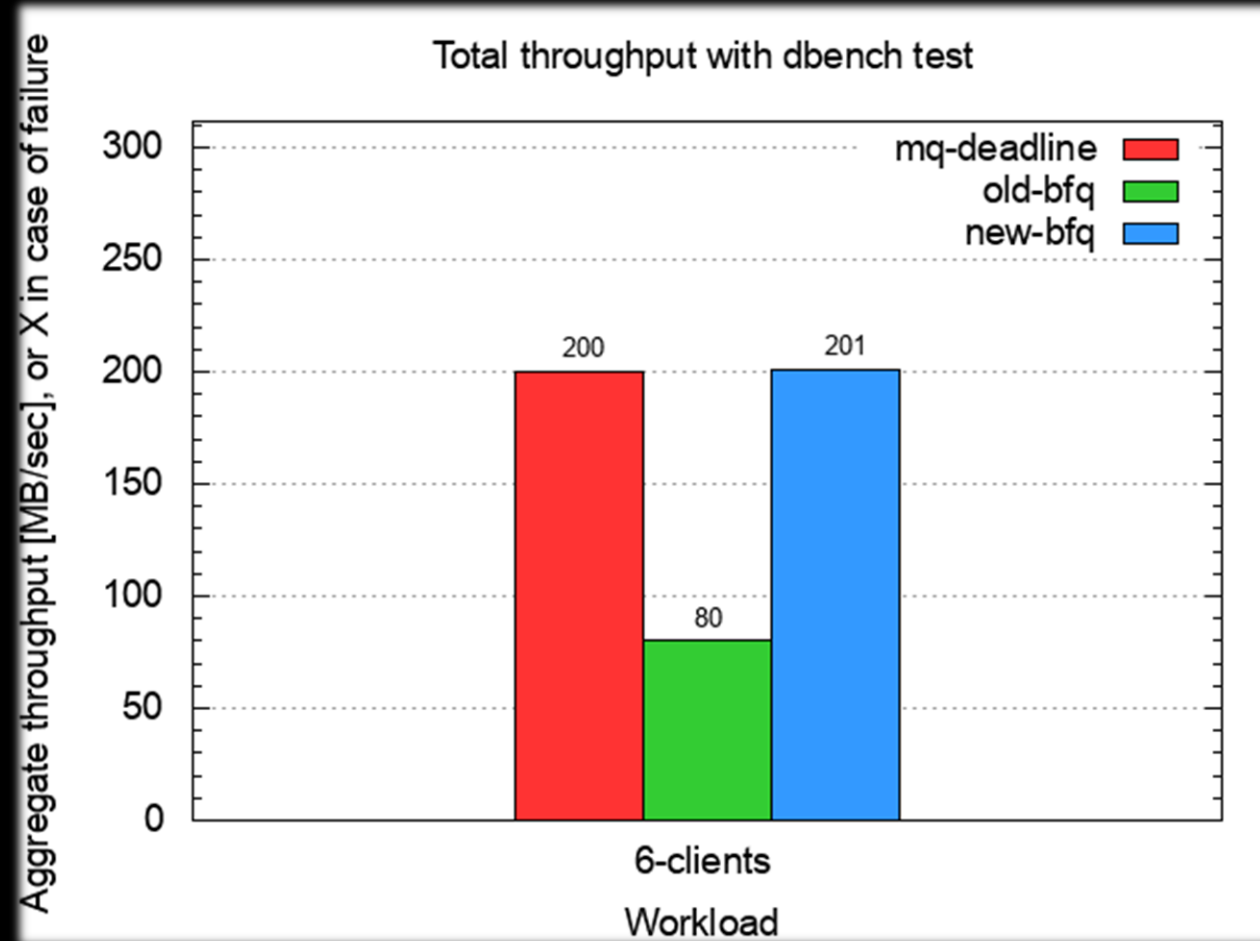


Background **BFQ**


Budget Fair Queueing



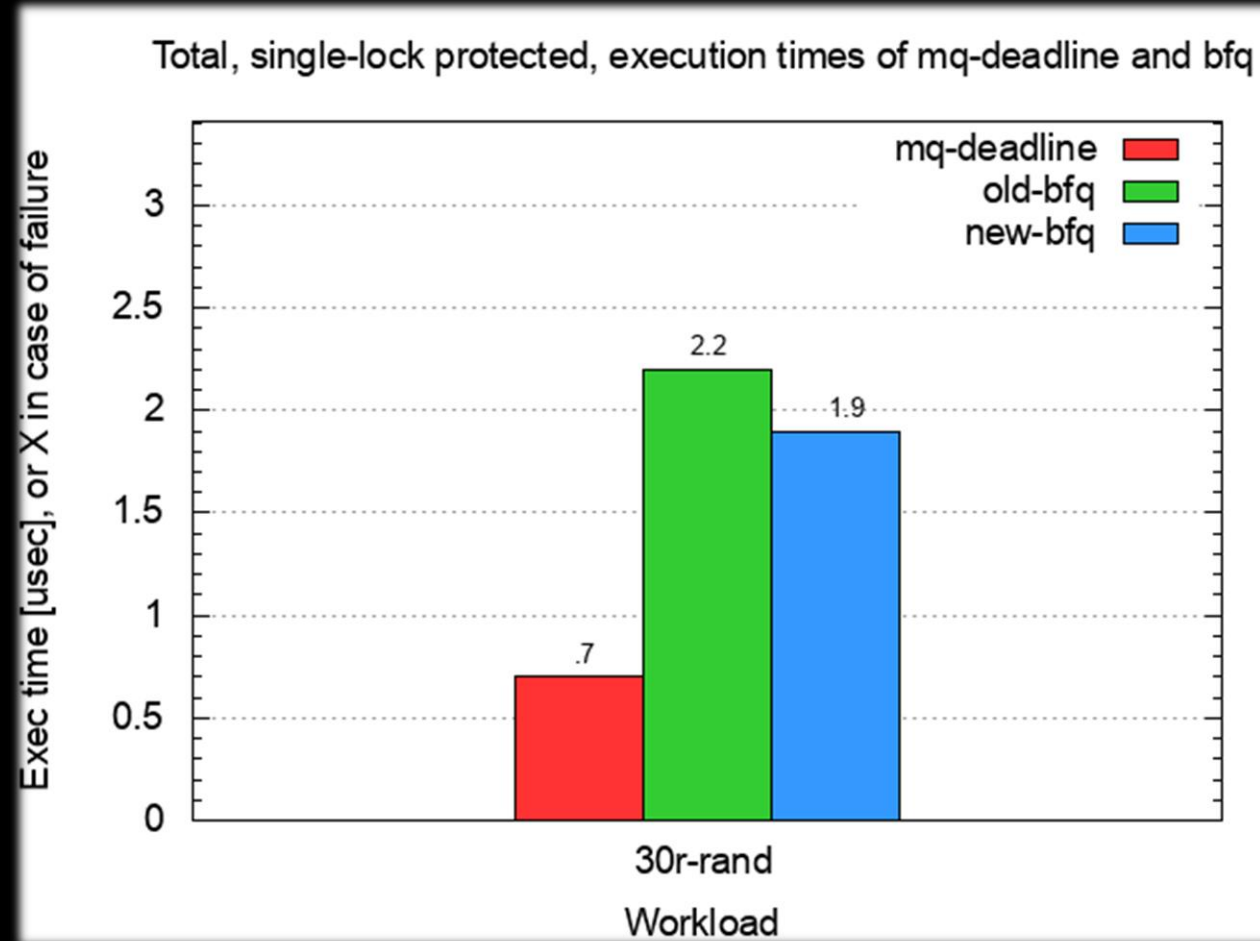
! BFQ improved



! BFQ improved



! BFQ improved



! BFQ improved

Tune service injection


Increase idling for weight-raised queues

Not merge queues on flash with queueing

Protect newly-created queues

! BFQ improved

Tune service injection



Increase idling for weight-raised queues



Not merge queues on flash with queueing



Protect newly-created queues

10


! BFQ improved

Tune service injection

Increase idling for weight-raised queues

Not merge queues on flash with queueing

Protect newly-created queues



10


! BFQ improved

Tune service injection

Increase idling for weight-raised queues

Not merge queues on flash with queueing

Protect newly-created queues



10

! BFQ improved

Tune service injection

Increase idling for weight-raised queues

Not merge queues on flash with queueing

Protect newly-created queues

merge



IO

merge




! BFQ improved

Tune service injection

Increase idling for weight-raised queues

Not merge queues on flash with queueing

Protect newly-created queues



10

5.3

+ pidfd_open(2)



+ pidfd_open(2)


```
pidfd = pidfd_open(pid);
```

```
check_if_pidfd_is_right_one(pidfd);
```

```
pidfd_send_signal(pidfd);
```


```
close(pidfd);
```

+ Clamping in scheduler




Background **EAS**


Energy Aware Scheduling





Background **EAS**



Background **EAS**




Background **EAS**



Background **Frequency control**



Background Frequency control




Background **Frequency control**




Background Frequency control




Background Frequency control




Background Frequency control




+ Clamping in scheduler




+ Clamping in scheduler




+ Clamping in scheduler



+ Clamping in scheduler



+ Clamping in scheduler




+ Clamping in scheduler

1. A **system call**, `sched_setattr()`
2. Two **cgroup knobs**, `util.{min, max}`
3. Two **sysctl knobs**, `sched_uclamp_util_{min, max}`

! Scan anons on file re-faults




Background LRU




Background Heuristic on LRU



Background Heuristic on LRU



Background Heuristic on LRU



Background **Re-fault**




Background **Re-fault**




Background **Re-fault**




! Scan anons on file re-faults




! Scan anons on file re-faults




! Scan anons on file re-faults




! Scan anons on file re-faults




! Scan anons on file re-faults



! Scan anons on file re-faults



! Scan anons on file re-faults



Q & A